

## Purely Functional Data Structures

Daniel Alonso

Introduction

Amortization

Fixing amortized analysis with laziness

Removing amortization

# Introduction


# What is “purely functional”?

Referential transparency: substituting an expression by its value does not change the program's behavior.

For data structures, it is essentially **immutability** (though there are other restrictions, e.g. not using randomization without explicitly passing a seed).

## Purely functional arrays

How do we modify a purely functional array? We don't. Think copy on write.




Cost of “modifying” a single array element:  $\Theta(n)$ .

# Lists




# Lists



When “modifying” an element, we must copy all parts of the data structure that point (directly or indirectly) to it.

# Lists




When “modifying” an element, we must copy all parts of the data structure that point (directly or indirectly) to it.


Cost of “modifying” a single list element:  $O(n)$ .



## Binary trees




## Binary trees




Insert 4 as a child of 3

## Binary trees




All nodes that point (directly or indirectly) to 3: path to the root.

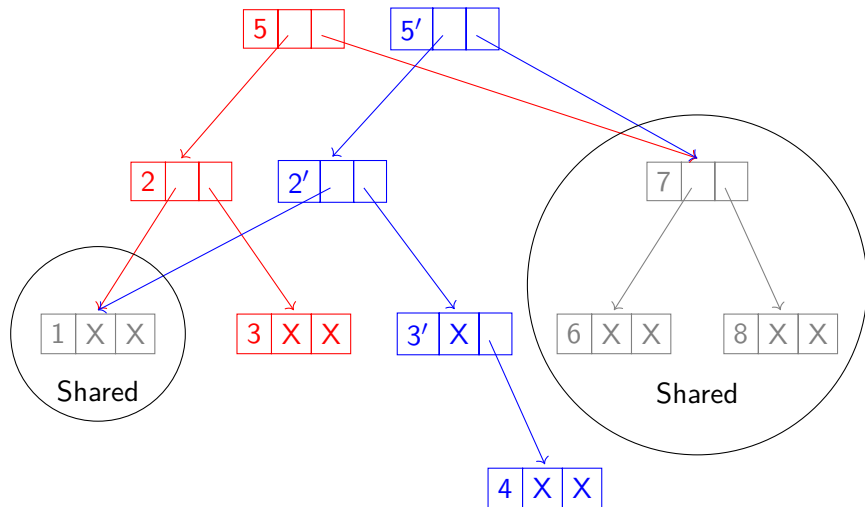
# Binary trees



# Binary trees



## Binary trees



In a balanced tree, we will copy  $O(\log n)$  nodes.

## Persistence

Because we copy (instead of modify) data structures, we keep both the input and the output of every operation. All the “versions” of a data structure that have ever existed in our program remain available (unless garbage collected).

## Running example: functional queue

```
class Queue q where
  empty :: q a -- trivial
  isEmpty :: q a -> Bool -- trivial
  head :: q a -> Maybe a
  snoc :: q a -> a -> q a -- snoc is rev of cons
  tail :: q a -> Maybe (q a)
```



## Simplest queue: list

```
data ListQueue a = LQ [a]
instance Queue ListQueue where
```

Function	Cost
head (LQ []) = Nothing	$\Theta(1)$
head (LQ (x:_)) = Just x	
snoc (LQ xs) x = LQ (xs ++ [x])	$\Theta(n)$
tail (LQ []) = Nothing	$\Theta(1)$
tail (LQ (_:xs)) = Just (LQ xs)	

## Batched queue: code

```
data BatchedQueue a = BQ [a] [a]
```

- Invariant: BQ fs rs is empty  $\iff$  fs is empty.
- rs contains the last elements of the queue in reverse order.

Example: we can represent 

1	2	3	4	5	6
---	---	---	---	---	---

 as BQ [1,2,3] [6,5,4].

## Batched queue: code

```
data BatchedQueue a = BQ [a] [a]
```

- Invariant:  $\text{BQ fs rs}$  is empty  $\iff$   $\text{fs}$  is empty.
- $\text{rs}$  contains the last elements of the queue in reverse order.

Example: we can represent 

1	2	3	4	5	6
---	---	---	---	---	---

 as  $\text{BQ } [1,2,3] \ [6,5,4]$ .

Function	Cost
<code>head (BQ [] _) = Nothing</code> <code>head (BQ (x:_) _) = Just x</code>	$\Theta(1)$
<code>snoc (BQ fs rs) x = check fs (x:rs)</code>	$\Theta(1)$
<code>tail (BQ [] _) = Nothing</code> <code>tail (BQ (_:fs) rs) = Just (check fs rs)</code>	$O(n)$
<code>check [] rs = BQ (reverse rs) []</code> <code>check fs rs = BQ fs rs</code>	

## Batched queue: example

BQ [] []

## Batched queue: example

```
snoc (BQ fs rs) x = check fs (x:rs)
check [] rs = BQ (reverse rs) []
check fs rs = BQ fs rs
```

BQ [] []

↓ snoc 1      (check [] [1] → fs empty, reverse rs)

BQ [1] []

## Batched queue: example

```
snoc (BQ fs rs) x = check fs (x:rs)
check [] rs = BQ (reverse rs) []
check fs rs = BQ fs rs
```

```
BQ [] []
  ↓ snoc 1   (check [] [1] → fs empty, reverse rs)
BQ [1] []
  ↓ snoc 2   (check [1] [2] → fs not empty)
BQ [1] [2]
```

## Batched queue: example

```
snoc (BQ fs rs) x = check fs (x:rs)
check [] rs = BQ (reverse rs) []
check fs rs = BQ fs rs
```

```
BQ [] []
  ↓ snoc 1   (check [] [1] → fs empty, reverse rs)
BQ [1] []
  ↓ snoc 2   (check [1] [2] → fs not empty)
BQ [1] [2]
  ↓ snoc 3   (same)
BQ [1] [3,2]
```

## Batched queue: example

```
tail (BQ (_:fs) rs) = Just (check fs rs)
check [] rs = BQ (reverse rs) []
check fs rs = BQ fs rs
```

BQ [] []

↓ snoc 1 (check [] [1] → fs empty, reverse rs)

BQ [1] []

↓ snoc 2 (check [1] [2] → fs not empty)

BQ [1] [2]

↓ snoc 3 (same)

BQ [1] [3,2]

$\Theta(n)$  ↓ tail (check [] [3,2] → fs empty, reverse rs)

Just (BQ [2,3] [])



## Batched queue: example

```
tail (BQ (_:fs) rs) = Just (check fs rs)
check [] rs = BQ (reverse rs) []
check fs rs = BQ fs rs
```

BQ [] []

↓ snoc 1 (check [] [1] → fs empty, reverse rs)

BQ [1] []

↓ snoc 2 (check [1] [2] → fs not empty)

BQ [1] [2]

↓ snoc 3 (same)

BQ [1] [3,2]

$\Theta(n)$  ↓ tail (check [] [3,2] → fs empty, reverse rs)

Just (BQ [2,3] [])

$\Theta(1)$  ↓ tail (check [3] [] → fs not empty)

Just (BQ [3] [])

# Amortization


## Amortized analysis

- Worst case of tail for BatchedQueue is  $\Theta(n)$ , though most tail operations will be  $\Theta(1)$ .
- Amortized analysis: bound the (worst case) cost of a sequence of operations, instead of each operation individually.

# Amortized analysis

- Worst case of tail for BatchedQueue is  $\Theta(n)$ , though most tail operations will be  $\Theta(1)$ .
- Amortized analysis: bound the (worst case) cost of a sequence of operations, instead of each operation individually.

Example:



Although there are  $n$  calls to `tail`, the cost of the whole sequence of operations is  $\Theta(n)$  and not  $\Theta(n^2)$ .

## Amortized analysis (ii)

For a sequence of  $m$  operations, define:

- $t_i$ : actual cost of operation  $i$ .
- $a_i$ : amortized cost of operation  $i$ .

We want to prove that:

$$\sum_{i=1}^m t_i \leq \sum_{i=1}^m a_i,$$

that is, the total actual cost is bounded above by the total amortized cost.

## Banker's method

- Each operation can deposit or spend some (imaginary) credits at different locations of the data structure.
  - Deposit: pay in advance to reduce the cost of a future operation.
  - Spend: take advantage of a “discount” paid for by a past operation.
- Define the amortized cost of operation  $i$  as:

$$a_i = t_i + c_i - \bar{c}_i, \quad \text{where } \begin{cases} c_i : \text{credits deposited} \\ \bar{c}_i : \text{credits spent} \end{cases}$$

- Credits can be spent **only once**.

## Banker's method

- Each operation can deposit or spend some (imaginary) credits at different locations of the data structure.
  - Deposit: pay in advance to reduce the cost of a future operation.
  - Spend: take advantage of a “discount” paid for by a past operation.
- Define the amortized cost of operation  $i$  as:

$$a_i = t_i + c_i - \bar{c}_i, \quad \text{where } \begin{cases} c_i : \text{credits deposited} \\ \bar{c}_i : \text{credits spent} \end{cases}$$

- Credits can be spent only once.

This is an analysis tool. These credits **do not exist** in code!


## Batched queue: amortized analysis

- head: does not deposit nor spend any credits,  $a_i = t_i = 1$




## Batched queue: amortized analysis

- head: does not deposit nor spend any credits,  $a_i = t_i = 1$
- snoc: empty          no credit deposited or spent,  $a_i = t_i = 1$




## Batched queue: amortized analysis

- head: does not deposit nor spend any credits,  $a_i = t_i = 1$
- snoc:   empty           no credit deposited or spent,  $a_i = t_i = 1$   
          not empty       deposit one credit to the new head of the rear list,  $a_i = t_i + 1 = 2$



## Batched queue: amortized analysis

- head: does not deposit nor spend any credits,  $a_i = t_i = 1$
- snoc:   empty           no credit deposited or spent,  $a_i = t_i = 1$   
          not empty       deposit one credit to the new head of the rear list,  $a_i = t_i + 1 = 2$
- tail:   no rotation   no credit deposited or spent,  $a_i = t_i = 1$




## Batched queue: amortized analysis

- head: does not deposit nor spend any credits,  $a_i = t_i = 1$
- snoc:   empty           no credit deposited or spent,  $a_i = t_i = 1$   
          not empty       deposit one credit to the new head of the rear list,  $a_i = t_i + 1 = 2$
- tail:   no rotation     no credit deposited or spent,  $a_i = t_i = 1$   
          rotation       spend all credits in the rear list,  $a_i = t_i - m = (1 + m) - m = 1$



## Batched queue: amortized analysis

- head: does not deposit nor spend any credits,  $a_i = t_i = 1$
- snoc:   empty           no credit deposited or spent,  $a_i = t_i = 1$   
          not empty       deposit one credit to the new head of the rear list,  $a_i = t_i + 1 = 2$
- tail:   no rotation   no credit deposited or spent,  $a_i = t_i = 1$   
          rotation       spend all credits in the rear list,  $a_i = t_i - m = (1 + m) - m = 1$




## Batched queue: amortized analysis


- head: does not deposit nor spend any credits,  $a_i = t_i = 1$
- snoc:   empty           no credit deposited or spent,  $a_i = t_i = 1$   
          not empty       deposit one credit to the new head of the rear list,  $a_i = t_i + 1 = 2$
- tail:   no rotation     no credit deposited or spent,  $a_i = t_i = 1$   
          rotation       spend all credits in the rear list,  $a_i = t_i - m = (1 + m) - m = 1$

Function	Worst case	Amortized
head	$\Theta(1)$	$\Theta(1)$
snoc	$\Theta(1)$	$\Theta(1)$
tail	$O(n)$	$\Theta(1)$

## Persistence breaks amortization




## Persistence breaks amortization






## Persistence breaks amortization



## Persistence breaks amortization



Persistence allows performing expensive operations an arbitrary number of times.  $n$  snocs followed by  $n$  tails on  $\boxed{1 \mid n \mid \dots \mid 2}$  have  $\Theta(n^2)$  cost. The amortized worst case cost for tail with persistence must be  $\Theta(n)$ .

## Things to remember

- Amortized analysis bounds the cost of a sequence of operations.
- In the banker's method, operations deposit (imaginary) credits to pay for the cost of other operations in advance.
- Persistence breaks standard amortized analysis.

Fixing amortized analysis with laziness


# Laziness reminder

Call-by-need:


- Do not evaluate arguments to functions before evaluating the function.

Call-by-name

`square x = x * x`



`square (prime 100)`



# Laziness reminder


Call-by-need:

- Do not evaluate arguments to functions before evaluating the function.
- Memoize the value of named expressions (difference with call-by-name).


`square x = x * x`

`square (prime 100)`

Call-by-name



Call-by-need



## Amortization with laziness

Divide the complete cost of operation  $i$  between:

**Unshared:** cost assuming every prior suspension is already memoized.

**Shared:** cost of forcing every suspension created but not evaluated by  $i$ .

Example:

$$\begin{aligned} & [] ++ ys = ys \\ & (x:xs) ++ ys = x : (xs ++ ys) \\ \\ & [1,2,3,4] ++ [5,6] \\ & \underbrace{1 : }_{\substack{\text{unshared} \\ (1)}} \underbrace{([2,3,4] ++ [5,6])}_{\substack{\text{shared} \\ (3)}} \end{aligned}$$

## Amortization with laziness

Divide the complete cost of operation  $i$  between:

**Unshared:** cost assuming every prior suspension is already memoized.

**Shared:** cost of forcing every suspension created but not evaluated by  $i$ .

Given a sequence of  $m$  operations, divide the shared costs between:

**Realized:** costs for suspensions evaluated at some point.

**Unrealized:** costs for suspensions never evaluated.

Example:

print \$ take 3 \$ [1,2,3,4] ++ [5,6]

↓

1	:	2	:	3	:	([4] ++ [5,6])
└──────────┘			└────────┘		└──────────────────────────┘	
unshared			realized		unrealized	
(1)			(2)		(1)	



## Amortization with laziness

Divide the complete cost of operation  $i$  between:

**Unshared:** cost assuming every prior suspension is already memoized.

**Shared:** cost of forcing every suspension created but not evaluated by  $i$ .

Given a sequence of  $m$  operations, divide the shared costs between:

**Realized:** costs for suspensions evaluated at some point.

**Unrealized:** costs for suspensions never evaluated.

$$\sum_{i=1}^m t_i = \sum_{i=1}^m (\text{unshared}_i + \text{realized}_i) \leq \sum_{i=1}^m a_i$$

## Adapting the Banker's method

Replace credits with **debits**: created to represent the debt of the *shared cost* of an operation. Must be paid before forcing a suspended computation.

$$a_i = \text{unshared}_i + d_i$$

$a_i$  takes into account debits discharged ( $d_i$ ), but not debits created (they could belong to unrealized suspensions).

## Adapting the Banker's method

Replace credits with **debits**: created to represent the debt of the *shared cost* of an operation. Must be paid before forcing a suspended computation.

$$a_i = \text{unshared}_i + d_i$$

$a_i$  takes into account debits discharged ( $d_i$ ), but not debits created (they could belong to unrealized suspensions).

Because no suspension is executed before paying its debits, we have:

$$\sum_{i=1}^m (\text{unshared}_i + \text{realized}_i) \leq \sum_{i=1}^m a_i$$

## Persistent queue: code

```
data PersistentQueue a = PQ Int [a] Int [a]
```

- In `PQ lenf fs lenr rs`, length of `fs`  $\geq$  length of `rs`.
- `rs` contains the last elements of the queue in reverse order.

Function	Cost
<pre>head (PQ _ [] _ _) = Nothing head (PQ _ (x:_) _ _) = Just x</pre>	$\Theta(1)$
<pre>snoc (PQ lenf fs lenr rs) x =     check lenf fs (lenr + 1) (x:rs)</pre>	$O(n)$
<pre>tail (PQ _ [] _ _) = Nothing tail (PQ lenf (_:fs) lenr rs) =     Just (check (lenf - 1) fs lenr rs)</pre>	$O(n)$
<pre>check lenf fs lenr rs       lenr &gt; lenf = PQ (lenf + lenr) (fs ++ reverse rs) 0 []       otherwise = PQ lenf fs lenr rs</pre>	


## Persistent queue: amortized analysis intuition

- First location in the queue must have no (unpaid) debit at any moment (so head can peek at it).
- All debits must have been discharged before a rotation.

## Persistent queue: amortized analysis intuition

- First location in the queue must have no (unpaid) debit at any moment (so head can peek at it).
- All debits must have been discharged before a rotation.

After a rotation, we create the following debits:




(fs becomes  $[0..n-1]$  ++ (reverse  $[2n..n]$ ))

## Persistent queue: amortized analysis intuition

- First location in the queue must have no (unpaid) debit at any moment (so head can peek at it).
- All debits must have been discharged before a rotation.


After a rotation, we create the following debits:



$\text{snoc } 2n$  discharges the first debit. Remaining debits:  $2n$ .

## Persistent queue: amortized analysis intuition

- First location in the queue must have no (unpaid) debit at any moment (so head can peek at it).
- All debits must have been discharged before a rotation.




After  $2n + 1$  snocs all debits must be discharged, as next snoc will cause a rotation. Each snoc discharges **1** debit.



## Persistent queue: amortized analysis intuition

- First location in the queue must have no (unpaid) debit at any moment (so head can peek at it).
- All debits must have been discharged before a rotation.



After  $n$  tails all debits must be discharged, as head must be able to peek at the first location. Each tail discharges 2 debits.


## Persistent queue: amortized analysis intuition

- First location in the queue must have no (unpaid) debit at any moment (so head can peek at it).
- All debits must have been discharged before a rotation.


$$a_i = \text{unshared}_i + d_i$$

Function	Worst case	unshared <sub><i>i</i></sub>	<i>d<sub>i</sub></i>	Amortized
head	$\Theta(1)$	$\Theta(1)$	0	$\Theta(1)$
snoc	$O(n)$	$\Theta(1)$	1	$\Theta(1)$
tail	$O(n)$	$\Theta(1)$	2	$\Theta(1)$

## Persistent queue: amortized bounds hold with persistence




## Persistent queue: amortized bounds hold with persistence



$[0..n-1] ++ (\text{reverse } [2n..n])$

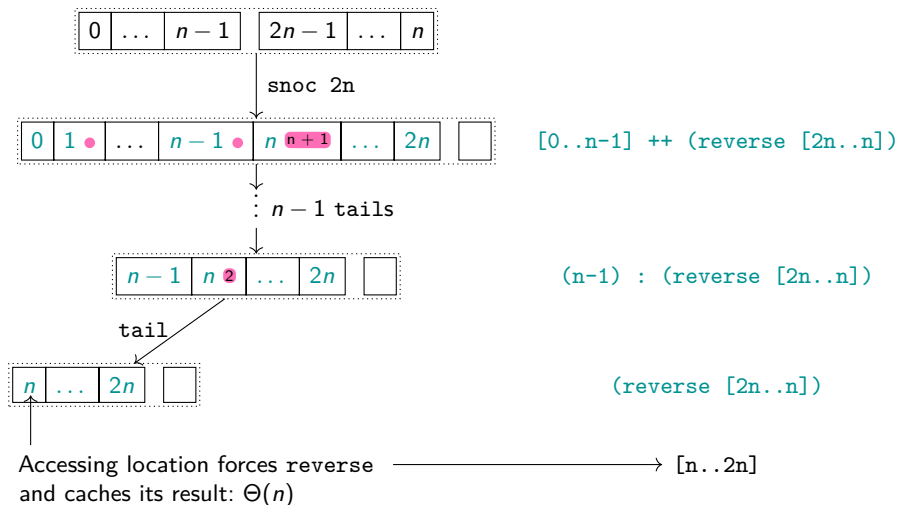
## Persistent queue: amortized bounds hold with persistence



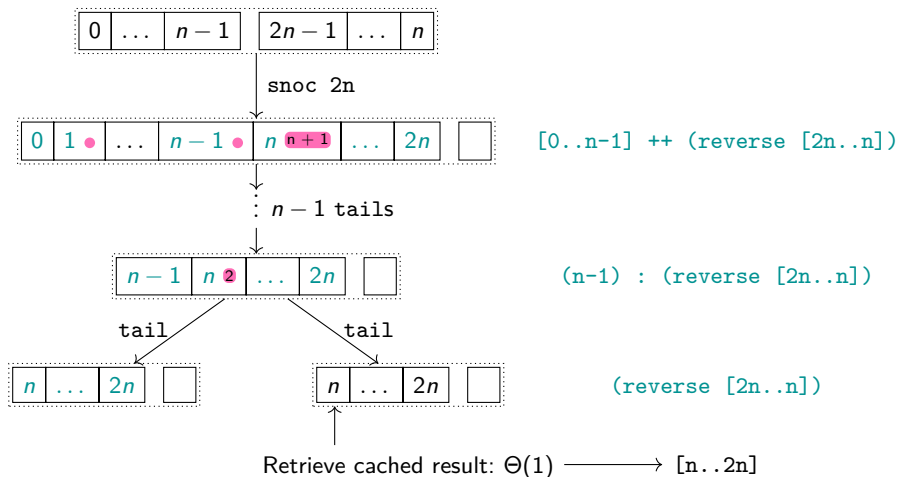
`[0..n-1] ++ (reverse [2n..n])`

`(n-1) : (reverse [2n..n])`

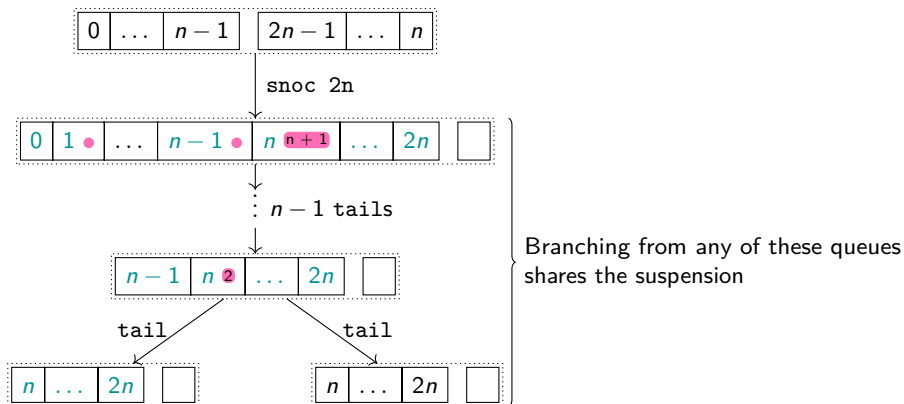
# Persistent queue: amortized bounds hold with persistence



# Persistent queue: amortized bounds hold with persistence

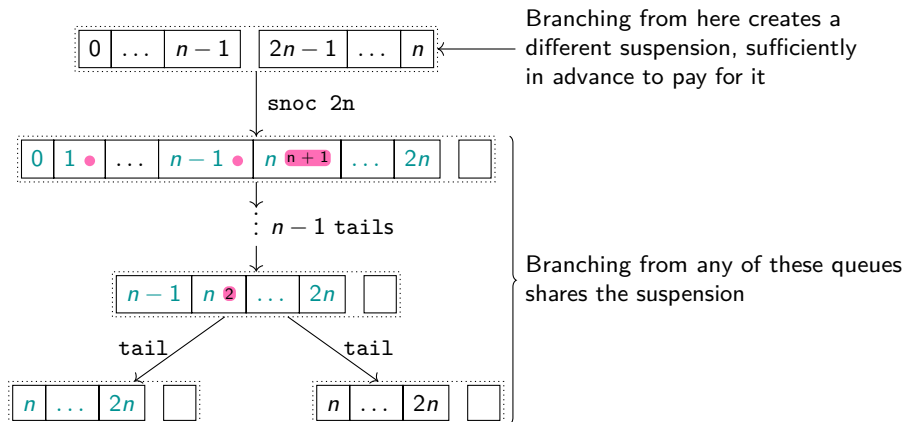


## Persistent queue: amortized bounds hold with persistence





# Persistent queue: amortized bounds hold with persistence



## Things to remember

- Credits are replaced by debits, representing the cost of suspended computations. Before forcing a suspension, all its debits must be discharged.
- Incremental suspended computations distribute its debits homogeneously, while monolithic ones accumulate all debits at the root.
- Monolithic suspended computations must be created sufficiently in advance. Their results will be shared when used persistently.

Removing amortization

## Why remove amortization?

Usually data structures with amortized bounds are simpler and/or faster, but sometimes bounding the time of every operation might be necessary:

- Interactive systems.
- Parallel systems.
- Real-time systems.

## Why remove amortization?

Usually data structures with amortized bounds are simpler and/or faster, but sometimes bounding the time of every operation might be necessary:

- Interactive systems.
- Parallel systems.
- Real-time systems.

Warning: non-strict languages like Haskell would usually require the use of `seq` to force specific suspensions, though we will ignore it.

# Scheduling

Similar idea to “strategies” from `parallel` Haskell package: decouple the creation of a value from the way it is evaluated.


Schedule: additional component of a data structure that forces pieces of a big suspension with every operation. It may require:

- Caution with chains of unevaluated suspensions.
- Transforming monolithic computations into incremental ones.

## Chains of unevaluated suspensions

```
list1 ++ list2  
(list1 ++ list2) ++ list3  
  ↓  
  ⋮  
  ↓  
(...((list1 ++ list2) ++ list3) ++ ...) ++ listk
```

## Chains of unevaluated suspensions



Although each suspension produced by `++` takes  $\Theta(1)$  time to force, we need to force  $k - 1$  suspensions to obtain the first element of the resulting list. Required time:  $\Theta(k)$ .



## From monolithic to incremental

When performing a rotation ( $fs \mathrel{++} reverse\ rs$ ), we know that  $|fs| + 1 = |rs|$ . Introduce an accumulating parameter (initially empty):

```
rotate    _      [r]    acc = y : acc
rotate (f:fs) (r:rs) acc = f : rotate fs rs (r : acc)
```

Idea: build the list from the beginning and from the end at the same time.

## From monolithic to incremental

When performing a rotation ( $fs \mathrel{++} reverse\ rs$ ), we know that  $|fs| + 1 = |rs|$ . Introduce an accumulating parameter (initially empty):

```
rotate    _      [r]    acc = y : acc
rotate (f:fs) (r:rs) acc = f : rotate fs rs (r : acc)
```

Idea: build the list from the beginning and from the end at the same time.

Example:

```
rotate [1,2,3] [7,6,5,4] []
1 : (rotate [2,3] [6,5,4] [7])
1 : 2 : (rotate [3] [5,4] [6,7])
1 : 2 : 3 : (rotate [] [4] [5,6,7])
1 : 2 : 3 : 4 : [5,6,7]
```

## Real-time queue


```
data RealTimeQueue a = RTQ [a] [a] [a]
```

In RTQ fs rs sch:

- sch contains the last elements of fs.
- $|sch| = |fs| - |rs|$ .


Function	Cost
<code>head (RTQ [] _ _) = Nothing</code>	$\Theta(1)$
<code>head (RTQ (x:_) _ _) = Just x</code>	
<code>snoc (RTQ fs rs sch) x = exec fs (x:rs) sch</code>	$\Theta(1)$
<code>tail (RTQ [] _ _) = Nothing</code>	$\Theta(1)$
<code>tail (RTQ (_:fs) rs sch) = Just (exec fs rs sch)</code>	
<code>exec fs rs (_:sch) = RTQ fs rs sch</code>	
<code>exec fs rs [] = RTQ fs' [] fs'</code> <code>where fs' = rotate fs rs []</code>	

# Real-time queue: example




## Real-time queue: example

```
snoc (RTQ fs rs sch) x = exec fs (x:rs) sch
exec fs rs (_:sch) = RTQ fs rs sch
exec fs rs [] = RTQ fs' [] fs'
  where fs' = rotate fs rs []
```




# Real-time queue: example

```
snoc (RTQ fs rs sch) x = exec fs (x:rs) sch
exec fs rs (_:sch) = RTQ fs rs sch
exec fs rs [] = RTQ fs' [] fs'
  where fs' = rotate fs rs []
```




# Real-time queue: example

```
snoc (RTQ fs rs sch) x = exec fs (x:rs) sch
exec fs rs (_:sch) = RTQ fs rs sch
exec fs rs [] = RTQ fs' [] fs'
  where fs' = rotate fs rs []
```



# Real-time queue: example


```
snoc (RTQ fs rs sch) x = exec fs (x:rs) sch
exec fs rs (_:sch) = RTQ fs rs sch
exec fs rs [] = RTQ fs' [] fs'
  where fs' = rotate fs rs []
```





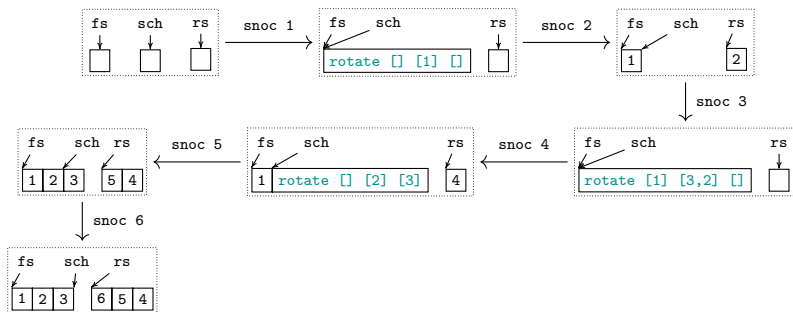
# Real-time queue: example

```
snoc (RTQ fs rs sch) x = exec fs (x:rs) sch
exec fs rs (_:sch) = RTQ fs rs sch
exec fs rs [] = RTQ fs' [] fs'
  where fs' = rotate fs rs []
```



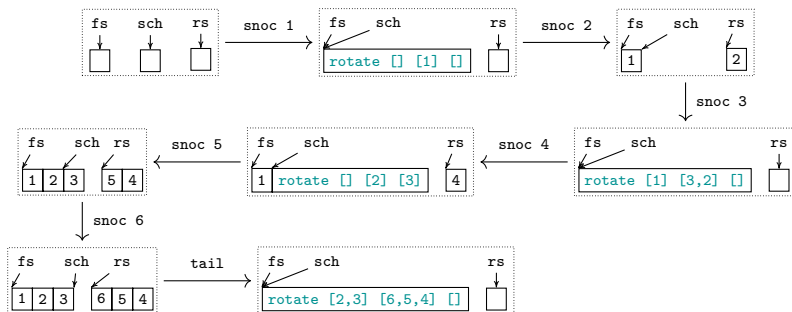
# Real-time queue: example

```
snoc (RTQ fs rs sch) x = exec fs (x:rs) sch
exec fs rs (_:sch) = RTQ fs rs sch
exec fs rs [] = RTQ fs' [] fs'
  where fs' = rotate fs rs []
```



# Real-time queue: example

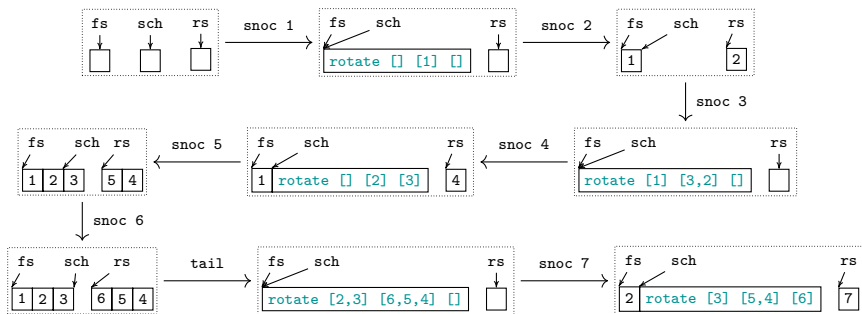
```
tail (RTQ (_:fs) rs sch) = Just (exec fs rs sch)
exec fs rs (_:sch) = RTQ fs rs sch
exec fs rs [] = RTQ fs' [] fs'
  where fs' = rotate fs rs []
```



# Real-time queue: example

```

snoc (RTQ fs rs sch) x = exec fs (x:rs) sch
exec fs rs (_:sch) = RTQ fs rs sch
exec fs rs [] = RTQ fs' [] fs'
    where fs' = rotate fs rs []
  
```



## Things to remember

- Avoid chains of unevaluated suspensions.
- Transform monolithic suspensions into incremental ones.
- Extend data structures with a schedule that will traverse them, forcing suspensions at convenient times.

Thanks!